



# A0/A1 cost ranking by hardware regime (never mixed within a panel)

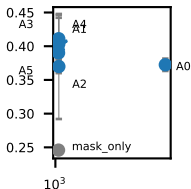
A: GPU-only



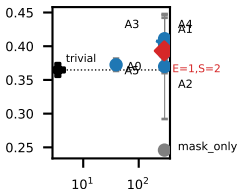
B1: CPU-only, 1 thread



B2: CPU-only, 60 threads\*



C: realistic (GPU model + CPU 60x-throughput det.)



latency / clip, ms (log scale)

\*B2 model: intra-op 60-thread call; B2 detection: 60-way multi-process throughput -- different parallelism mechanisms, both spending a 60-CPU-thread budget (see outputs/meld\_phase6\_regimes.json).

test macro-F1 (mean, 5 seeds)