


Model-only vs. end-to-end latency: the ranking reverses



Dotted segments join each condition's model-only and end-to-end latency.
 A1 end-to-end (detection included) costs $\sim 161\times$ its own model-only latency.